

# Knowledge Base Usage Report

**Project Title:** City Info Retrieval API

**Purpose:** To provide intelligent answers to user queries based on a city knowledge base using a Retrieval-Augmented Generation (RAG) system.

## ❖ Data Flow and Processing Overview

### 1. Knowledge Base Loading

- **Source:** JSON file (JSON\_FILE\_PATH) containing city-related information.
- **Step:** `load_and_prepare_vectorstore()`
- **Action:** Converts city data into vector embeddings for semantic search.

### 2. RAG Chain Setup

- **Step:** `setup_rag_chain(vectorstore)`
- **Action:** Links a retriever (vectorstore) with a language model to enable question-answering based on retrieved content.

### 3. Handling User Queries

- **Endpoint:** POST /query
- **Model:** QueryRequest
- **Action:**
  - Accepts a user question.
  - Checks for cached results.
  - If not cached:
    - Retrieves relevant content using the vectorstore.
    - Generates a response via the RAG chain.
    - Caches the response.

### 4. Searching the Knowledge Base

- **Endpoint:** POST /search
- **Action:** Returns the top-k semantically relevant entries from the vectorstore based on user input.

### 5. Accessing Raw Knowledge Base

- **Endpoint:** GET /data

- **Action:** Allows users to view all original knowledge base entries

## **Tools Implemented**

- FastAPI: Web framework for API routes
- Pydantic: Validates data models
- LangChain and RAG: Handles retrieval and generation
- Vectorstore: Enables semantic document search
- AsyncIO + ThreadPoolExecutor: Improves responsiveness
- Hashlib + LRU Cache: Speeds up repeated query results